



US006041345A

United States Patent [19]  
Levi et al.

[11] Patent Number: 6,041,345  
[45] Date of Patent: Mar. 21, 2000

[54] ACTIVE STREAM FORMAT FOR HOLDING  
MULTIPLE MEDIA STREAMS

07245600 9/1995 Japan ..... H04L 1/00

OTHER PUBLICATIONS

[75] Inventors: Steven P. Levi, Redmond; Mark D. VanAntwerp, Seattle; Craig M. Dowell, Redmond; Chadd B. Knowlton, Bellevue, all of Wash.

Huang, J.-H., et al., "MHTP—A Multimedia High-Speed Transport Protocol", *Proceedings, IEEE GLOBECOM '92*, vol. 3, p. 1364–1368 (Dec. 6–9, 1992).

LaPorta, T.F., et al., "The MultiStream Protocol: A Highly Flexible High-Speed Transport Protocol", *IEEE Journal on Selected Areas in Communications*, 11, p. 519–530 (May 1993).

[73] Assignee: Microsoft Corporation, Redmond, Wash.

Sarginson, P.A., "MPEG–2: A Tutorial Introduction to the Systems Layer", *IEE Colloquium on MPEG—What It Is and What It Isn't*, p. 4/1–4/13, (Jan. 1, 1995).

[21] Appl. No.: 08/813,151

Schatzmayr, R., et al., "Providing Support for Data Transfer in a New Networking Environment", *Multimedia Transport and Teleservices. International Cost 237 Works Proceedings, Vienna, Austria*, p. 241–255, (Nov. 13, 1994).

[22] Filed: Mar. 7, 1997

Ohta, N., *Packet Video: Modeling and Signal Processing*, Norwood, MA: Artech House, Inc., 144–153, (1994).

Related U.S. Application Data

[60] Provisional application No. 60/013,029, Mar. 8, 1996, and provisional application No. 60/028,789, Oct. 21, 1996.

[51] Int. Cl.<sup>7</sup> ..... G06F 17/50

[52] U.S. Cl. .... 709/217; 445/34.1

[58] Field of Search ..... 348/38, 441; 370/232, 370/17, 94.3, 94.1, 312, 474, 329; 375/46; 371/37.4, 40; 369/275.3; 386/111; 395/325; 709/217; 380/14; 711/202; 455/4.2, 34.1

Primary Examiner—Frank J. Asta  
Assistant Examiner—Thong Vu  
Attorney, Agent, or Firm—Schwegman, Lundberg, Woessner & Kluth, P.A.

[56] References Cited

[57] ABSTRACT

U.S. PATENT DOCUMENTS


5,319,707	6/1994	Wasileski et al. ....	380/14
5,467,342	11/1995	Logston et al. ....	370/17
5,506,847	4/1996	Shobatake ....	370/94.3
5,559,813	9/1996	Shimizu ....	371/37.4
5,604,843	2/1997	Shaw et al. ....	395/101
5,625,877	4/1997	Dunn et al. ....	455/34.1
5,654,962	8/1997	Rostoker et al. ....	370/232
5,668,803	9/1997	Tymes et al. ....	370/312
5,671,226	9/1997	Murakani et al. ....	370/474
5,708,961	1/1998	Hylton et al. ....	455/4.2
5,754,242	5/1998	Ohkami ....	348/441
5,754,589	5/1998	Maitra et al. ....	375/216
5,774,461	6/1998	Hyden et al. ....	370/329
5,774,481	6/1998	Hyden et al. ....	370/329
5,842,224	11/1998	Fenner ....	711/202
5,911,776	6/1999	Guck ....	709/217

An active stream format is defined and adopted for a logical structure that encapsulates multiple data streams. The data streams may be of different media. The data of the data streams is partitioned into packets that are suitable for transmission over a transport medium. The packets may include error correcting information. The packets may also include clock licenses for dictating the advancement of a clock when the data streams are rendered. The format of ASF facilitates flexibility and choice of packet size and in specifying maximum bit rate at which data may be rendered. Error concealment strategies may be employed in the packetization of data to distribute portions of samples to multiple packets. Property information may be replicated and stored in separate packets to enhance its error tolerance. The format facilitates dynamic definition of media types and the packetization of data in such dynamically defined data types within the format.

FOREIGN PATENT DOCUMENTS

0 753 954 A2 1/1997 European Pat. Off. .... H04L 29/06

21 Claims, 24 Drawing Sheets




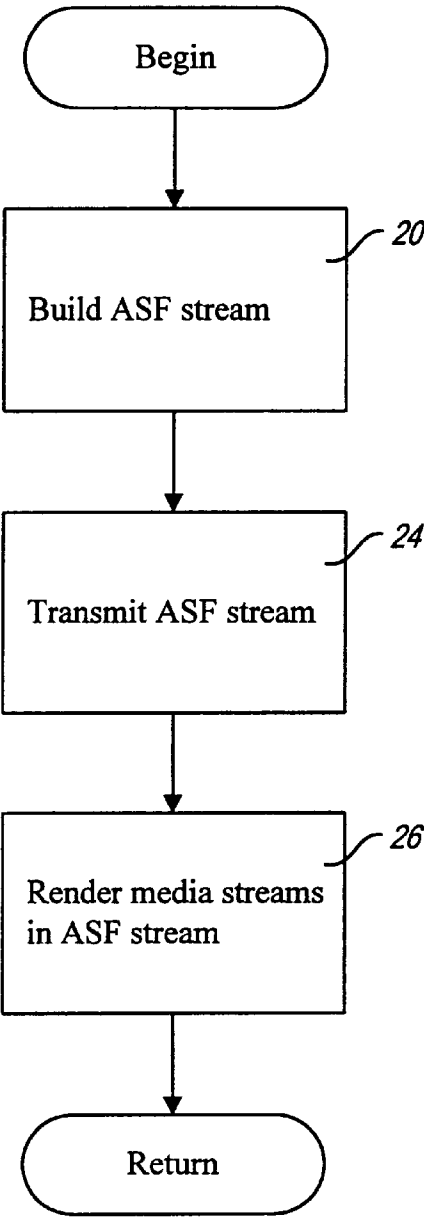


Fig. 1



*Fig. 2*

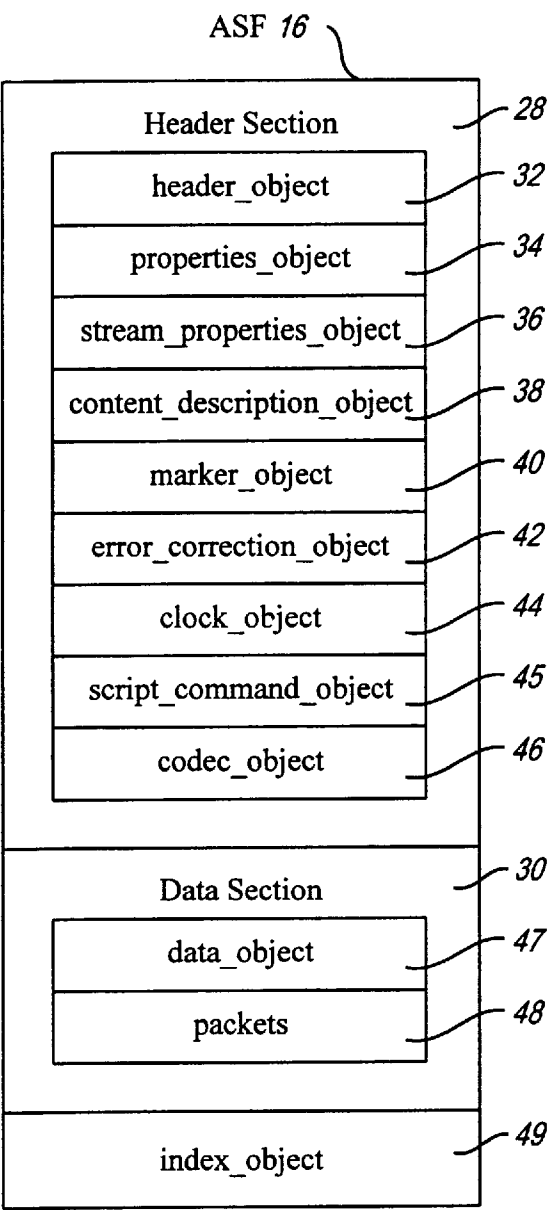


Fig. 3




Fig. 4





Fig. 5

*Fig. 6A*


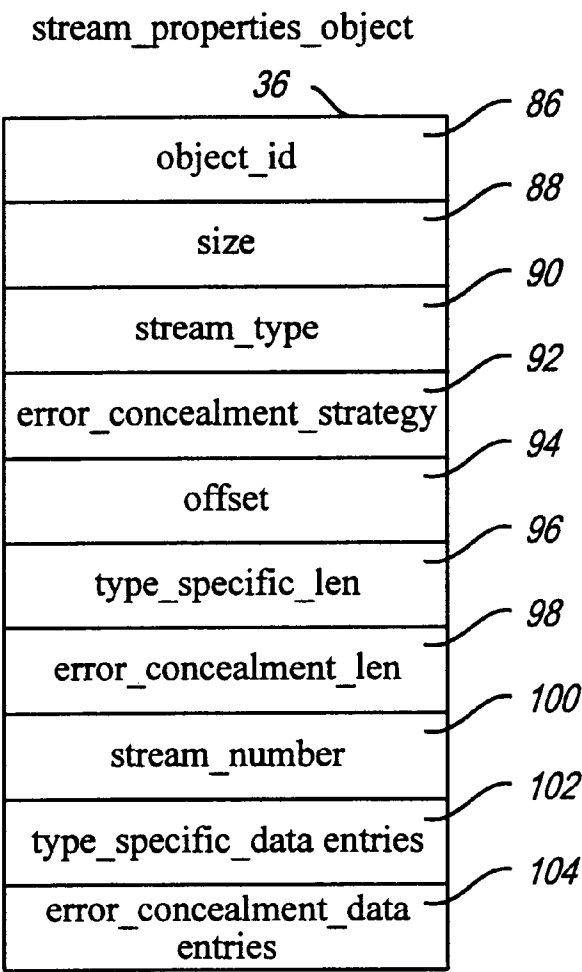
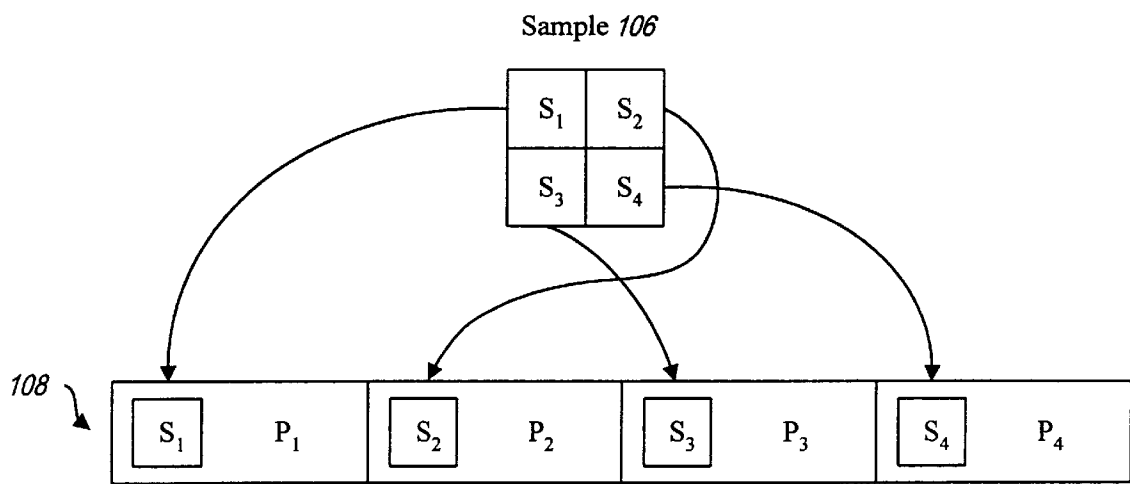


Fig. 6B



*Fig. 7*





*Fig. 8*






Fig. 9



*Fig. 10A*



*Fig. 10B*

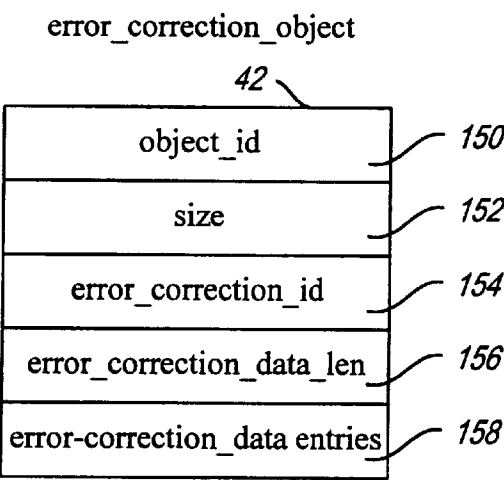





Fig. 11

*Fig. 12*



*Fig. 13*



*Fig. 14A*






Fig. 14B





Fig. 14C



*Fig. 15A*





Fig. 15B



*Fig. 16*

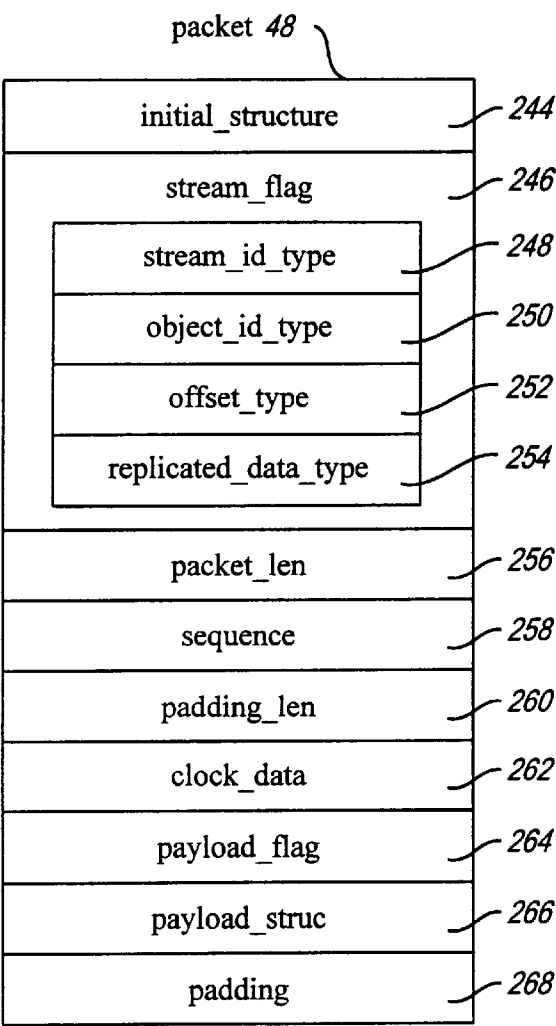

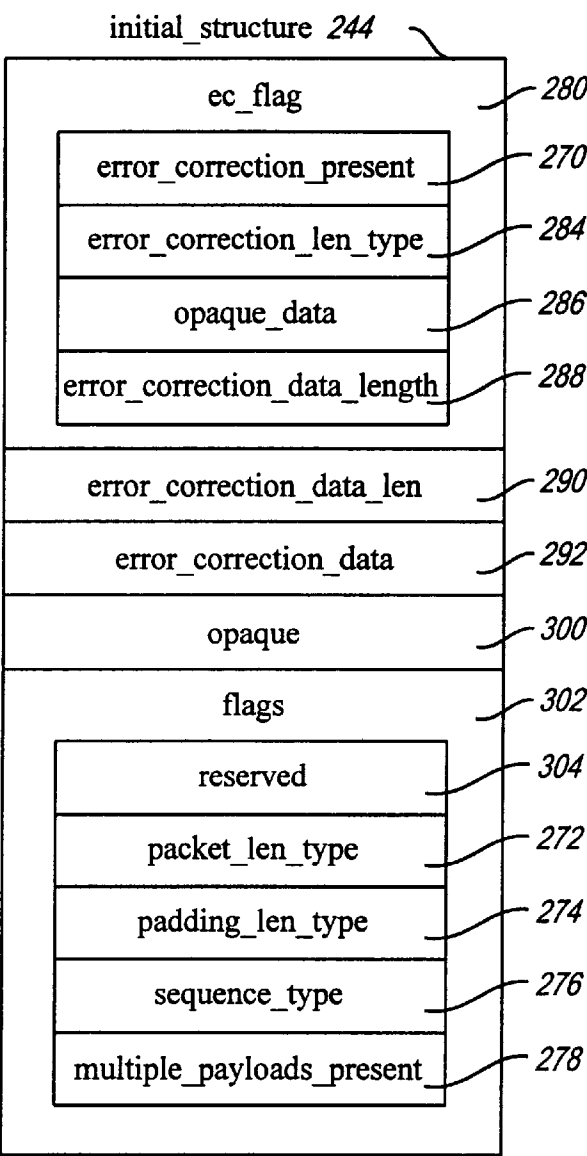



Fig. 17



*Fig. 18A*



*Fig. 18B*



*Fig. 19*






Fig. 20

## ACTIVE STREAM FORMAT FOR HOLDING MULTIPLE MEDIA STREAMS

This application claims benefit of Provisional Application Ser. No. 60/013,029 filed Mar. 8, 1996, and a provisional of Ser. No. 60/028,789 filed Oct. 21, 1996.

### TECHNICAL FIELD

The present invention relates generally to data processing systems and more particularly to an active stream format for holding multiple media streams.

### BACKGROUND OF THE INVENTION

Conventional file and/or stream formats for transmitting multiple data streams of varying media are limited in several respects. First, these formats are generally limited in the packet sizes that are available for encapsulating data. Such formats, if they specify packets, specify the packets as a given fixed size. Another limitation of such formats is that they do not facilitate the use of error correction codes. A further weakness of these conventional formats is that they do not provide flexibility in timing models for rendering the data encapsulated within the format. An additional limitation with such formats is that they are not well adapted for different transport mediums that have different levels of reliability and different transmission capabilities.

### SUMMARY OF THE INVENTION

In accordance with a first aspect of the present invention, a computer system has a logical structure for encapsulating multiple streams of data that are partitioned into packets for holding samples of data from the multiple data streams. A method of incorporating error correction into the logical structure is performed on the computer system. In accordance with this method, a portion of at least one packet is designated for holding error correcting data. The error correcting data is then stored in the designated portion of the packet.

In accordance with another aspect of the present invention, multiple streams of data are stored in packets and error correcting data is stored in at least some of the packets. The packets are encapsulated into a larger stream and information regarding what error correcting methods are employed for the packets is also stored in the packets.

In accordance with yet another aspect of the present invention, samples of data from multiple data streams are stored in packets, and replicas of information are stored in at least some of the packets. A flag is set in each of the packets that holds replicas to indicate that the packets hold the replicas. The packets are encapsulated into a larger logical structure and transmitted to a destination.

In accordance with a further aspect of the present invention, a logical structure is provided for encapsulating multiple streams of data where the streams of data are stored in packets. Clock licenses that dictate advancement of a clock are stored in multiple ones of the packets. The logical structure is transmitted from a source computer to a destination computer. The clock is advanced at the destination computer as dictated by the clock license for each packet that holds a clock license in response to the receipt or processing of the packet at the destination computer.

In accordance with an additional aspect of the present invention, a stream format is provided for encapsulating multiple streams of data. The stream format includes a field for specifying a packet size for holding samples of the

multiple streams of data. In a logical structure that adopts the stream format, a value is stored in the field that corresponds to the desired packet size. Packets of the desired size are stored within the logical structure and the logical structure is transmitted over a transport medium to the destination.

In accordance with a further aspect of the present invention, a stream format is provided for encapsulating multiple streams of data. A field is included in a logical structure that adopts the stream format for holding a value that specifies a maximum bit rate at which the multiple streams may be rendered at the destination. A value is stored in the field and the logical structure is transmitted over a transport medium to a destination.

In accordance with another aspect of the present invention, a stream format is provided for encapsulating multiple data streams and a new media type is dynamically defined. An identifier of the media type is stored in a logical structure that adopts the stream format and packets of the new media type are stored in the logical structure.

### BRIEF DESCRIPTION OF THE DRAWINGS

A preferred embodiment of the present invention will be described below relative to the following drawings.

FIG. 1 is a block diagram illustrating a computer system that is suitable for practicing the preferred embodiment of the present invention.

FIG. 2 is a flowchart illustrating use of the ASF stream in accordance with a preferred embodiment of the present invention.

FIG. 3 is a block diagram illustrating the components of the ASF stream.

FIG. 4 is a block diagram illustrating the format of the header\_object.

FIG. 5 is a block diagram illustrating the format of the properties\_object.

FIG. 6A is a flowchart illustrating the steps that are performed to fill in packet size fields within the ASF stream.

FIG. 6B is a diagram illustrating different packet sizes and respective ASF streams.

FIG. 7 is a block diagram illustrating the format of the stream\_properties\_object.

FIG. 8 is a diagram that illustrates the partitioning of a sample for storage in multiple packets.

FIG. 9 is a diagram that illustrates the format of the content\_description\_object.

FIG. 10A is a diagram illustrating the format of the marker\_object.

FIG. 10B is a diagram illustrating the format of a marker entry.

FIG. 11 is a diagram illustrating the format of the error\_correction\_object.

FIG. 12 is flowchart illustrating the steps that are performed to utilize error correcting information in accordance with a preferred embodiment of the present invention.

FIG. 13 is a diagram illustrating format of the clock\_object.

FIG. 14A is a diagram illustrating the format of the script\_command\_object.

FIG. 14B is a diagram illustrating the format of a type\_names\_struct.

FIG. 14C is a diagram illustrating the format of a command\_entry.

FIG. 15A is a diagram illustrating the format of the codec\_object.

FIG. 15B is a diagram of a CodecEntry.

FIG. 16 is a diagram illustrating the format of the data\_\_object.

FIG. 17 illustrates the format of a packet.

FIG. 18A illustrates a first format that the initial\_\_structure may assume.

FIG. 18B illustrates a second format that the initial\_\_structure may assume.

FIG. 19 illustrates the format of a payload\_\_struc.

FIG. 20 is a diagram illustrating the format of the index\_\_object.

### DETAILED DESCRIPTION OF THE INVENTION

The preferred embodiment of the present invention employs an active stream format (ASF) for holding multiple media streams. ASF is well suited for storage of multimedia streams as well as transmission of multiple media streams over a transport medium. ASF is constructed to encapsulate diverse multimedia streams and facilitates optimal interleaving of respective media streams. ASF specifies the packetization of data and provides flexibility in choosing packet sizes. In addition, ASF enables the specification of a maximum data transmission rate. As such, the packetization and transmission of media streams may be tailored to facilitate the bandwidth limitations of the system on which media streams are stored or transmitted.

ASF facilitates the use of error correction and error concealment techniques on the media streams. In unreliable transport mediums, such error correction and error concealment is highly beneficial. ASF is independent of media types and is extensible to handle newly defined media types. ASF supports flexible timing approaches and allows an author of an ASF stream to specify the synchronization of events. ASF supports synchronized rendering using a variety of synchronization clock types and provides index information which can be used as markers for lookup to provide playback features such as fast forward and fast reverse.

FIG. 1 is a block diagram of an illustrative system for practicing the preferred embodiment of the present invention. FIG. 2 is a flowchart that illustrates the steps that are performed in the illustrative embodiment of FIG. 1. An ASF stream 16 is built by an author (step 20 in FIG. 2) and stored on a storage 14 on a source computer 10. As will be described in more detail below, ASF allows the author to design the stream for a most efficient storage based on the type of source computer 10 on which it is stored. Sometime later, the ASF stream 16 is transferred over a transport media 17, such as a network connection, to a destination computer 12 (step 24 in FIG. 2). The destination computer 12 includes a number of renderers 18 for rendering the media types that are present within the ASF stream 16. For example, the ASF stream 16 may include audio-type data and video-type data. The renderers 18 at the destination 12 include an audio renderer and a video renderer. The renderers may begin rendering data as soon as they receive data prior to the complete transmission of the entire ASF stream 16 (see step 26 in FIG. 2). The renderers need not immediately render the data, but rather may render the data at a later point in time.

FIG. 3 depicts the basic logical organization of an ASF stream 16. It is up to the author to fill in the contents of the ASF stream in accordance with this format. The ASF stream 16 is divisible into a header section 28, a data section 30 and an index section 49. In general, the header section is first transmitted from the source computer 10 to the destination

computer 12 so that the destination computer may process the information within the header section. Subsequently, the data section 30 is transmitted from the source computer 10 to the destination computer 12 on a packet-by-packet basis and the index section 49 is transmitted. The header section 28 includes a number of objects that describe the ASF stream 16 in aggregate. The header section 28 includes a header\_\_object 32 that identifies the beginning of the ASF header section 28 and specifies the number of objects contained within the header section. FIG. 4 depicts the format of the header\_\_object 32 in more detail. The header\_\_object 32 includes an object\_\_id field 50 that holds a UUID for the header\_\_object. The UUID is an identifier. The header\_\_object 32 also includes a size field 52 that specifies a 64-bit quantity that describes the size of the header section 28 in bytes. The header\_\_object 32 additionally includes a number\_\_headers field 54 that holds a 32-bit number that specifies a count of the objects contained within the header section that follow the header\_\_object 32. An alignment field 55 specifies packing alignment of objects within the header (e.g., byte alignment or word alignment). The architecture field 57 identifies the computer architecture type of the data section 30 at the index section 49. The architecture field 57 specifies the architecture of these sections as little endian or big endian.

The header\_\_object 32 is followed in the header section 28 by a properties\_\_object 34, such as depicted in FIG. 5. The properties\_\_object 34 describes properties about the ASF stream 16. As can be seen in FIG. 5, the properties\_\_object 34 includes an object\_\_id field 56 that holds a UUID and a size field 58 that specifies the size of the properties\_\_object 34. The properties\_\_object 34 also includes a multimedia\_\_stream\_\_id field 60 that contains a UUID that identifies a multimedia ASF stream. A total\_\_size field 62 is included in the properties\_\_object 34 to hold a 64-bit value that expresses the size of the entire ASF multimedia stream.

The properties\_\_object 34 also holds a created field 64 that holds a timestamp that specifies when the ASF stream was created. A num\_\_packet field 65 holds a 64-bit value that defines the number of packets in the data section 30. A play\_\_duration field 66 holds a 32-bit number that specifies the play duration of the entire ASF stream in 100-nanosecond units. For example, if the ASF stream 16 holds a movie, the duration field 66 may hold the duration of the movie. The play\_\_duration field 66 is followed by a send\_\_duration field 67 that corresponds to send the ASF stream in 100-nanosecond units. A preroll field 68 specifies the amount of time to buffer data before starting to play, and the flags field 70 holds 32-bits of bit flags.

The properties\_\_object 34 includes a min\_\_packet\_\_size field 72 and a max\_\_packet\_\_size field 74. These fields 72 and 74 specify the size of the smallest and largest packets 48 in the data section 30, respectively. These fields help to determine if the ASF stream 16 is playable from servers that are constrained by packet size. For constant bit rate streams, these values are set to have the same values. A maximum\_\_bit\_\_rate field 76 holds a value that specifies the maximum instantaneous bit rate (in bits per second) of the ASF stream.

FIG. 6A is a flowchart illustrating how these values are identified and assigned during authoring of the ASF stream 16. First, the size of the smallest packet in the data section 30 is identified (step 78 in FIG. 6A). The size of the smallest packet is stored in the min\_\_packet\_\_size field 72 (step 80 in FIG. 6A). The size of the largest packet in the data section 30 is identified (step 82 in FIG. 6A), and the size is assigned to the max\_\_packet\_\_size field 74 (step 84 in FIG. 6A).

One of the beneficial features of ASF is its ability for facilitating different packet sizes for data of multiple media

streams. FIG. 6B shows one example of two different streams **83** and **85**. In stream **83**, each of the packets is chosen to have a size of 512 bytes, whereas in stream **85** each of the packets **48** holds 256 bytes. The decision as to the size of the packets may be influenced by the speed of the transport mechanism over which the ASF stream is to be transmitted, the protocol adopted by the transport medium, and the reliability of the transport medium.

As mentioned above, the `properties_object` **34** holds a value in the `maximum_bit_rate` field **76** that specifies an instantaneous maximum bit rate in bits per second that is required to play the ASF stream **16**. The inclusion of this field **76** helps to identify the requirements necessary to play the ASF stream **16**.

The header section **28** (FIG. 3) must also include at least one `stream_properties_object` **36**. The `stream_properties_object` **36** is associated with a particular type of media stream that is encapsulated within the ASF stream **16**. For example, one of the `stream_properties_objects` **36** in the header section **28** may be associated with an audio stream, while another such object is associated with a video stream. FIG. 7 depicts a format for such `stream_properties_objects` **36**. Each `stream_properties_object` **36** includes an `object_id` field **86** for holding a UUID for the object and a size field **88** for holding a value that specifies the size of the object in bytes. A `stream_type` field **90** holds a value that identifies the media type of the associated stream.

The `stream_properties_object` **36** holds at least three fields **92**, **98** and **104** for holding information relating to error concealment strategies. In general, ASF facilitates the use of error concealment strategies that seek to reduce the effect of losing information regarding a given sample of media data. An example of an error concealment strategy is depicted in FIG. 8. A sample **106** is divided into four sections  $S_1$ ,  $S_2$ ,  $S_3$  and  $S_4$ . When the sample is incorporated into packets in the ASF stream, the samples are distributed into separate packets  $P_1$ ,  $P_2$ ,  $P_3$  and  $P_4$  so that if any of the packets are lost, the amount of data that is lost relative to the sample is not as great, and techniques, such as interpolation, may be applied to conceal the error. Each sample has a number of associated properties that describe how big the sample is, how the sample should be presented to a viewer, and what the sample holds. Since the loss of the property information could prevent the reconstruction of the sample, the properties information for the entire sample is incorporated with the portions of the sample in the packets.

The `error_concealment_strategy` field **92** holds a UUID that identifies the error concealment strategy that is employed by the associated stream. The `error_concealment_len` field **98** describes the number of bytes in an error concealment data block that is held in the `error_concealment_data` entries **104**. The properties associated with the error concealment strategy are placed in the `error_concealment_data` entries **104**. The number of entries will vary depending upon the error concealment strategy that is adopted.

The `stream_properties_object` **36** includes a `stream_number` field **100** that holds an alias to a stream instance. The `stream_properties_object` **36** also includes an `offset` field **94** that holds an offset value to the stream in milliseconds. This value is added to all of the timestamps of the samples in the associated stream to account for the offset of the stream with respect to the timeline of the program that renders the stream. Lastly, the `stream_properties_object` **36** holds a `type_specific_len` field **96** that holds a value that describes the number of bytes in the `type_specific_data` entries **102**.

The `type_specific_data` entries **102** hold properties values that are associated with the stream type.

The header section **28** (FIG. 3) may also include a number of optional objects **38**, **40**, **42**, **44**, **45** and **46**. These optional objects include a `content_description_object` **38** that holds information such as the title, author, copyright information, and ratings information regarding the ASF stream. This information may be useful and necessary in instances wherein the ASF stream **16** is a movie or other artistic work. The `content_description_object` **38** includes an `object_id` field **110** and a size field **112** like the other objects in the header section **28**. A `title_len` field **114** specifies the size in bytes of the title entries **119** that hold character data for the title of the ASF stream **16**. An `author_len` field **115** specifies the size in bytes of the author entries **120** which hold the characters that specify the author of the ASF stream **16**. The `copyright_len` field **116** holds the value that specifies the length in bytes of the copyright entries **121** that hold copyright information regarding the ASF stream **16**. The `description_len` field **117** holds a value that specifies the length in bytes of the description entries **122**. The description entries **122** hold a narrative description of the ASF stream **16**. Lastly, the `rating_len` field **118** specifies a size in bytes of the rating entries **123** that hold rating information (e.g., X, R, PG-13) for the ASF stream content.

The header section **28** may include a `marker_object` **40**. The `marker_object` **40** holds a pointer to a specific time within the data section **30**. The `marker_object` enables a user to quickly jump forward or backward to specific data points (e.g., audio tracks) that are designated by markers held within the `marker_object` **40**.

FIG. 10A shows the `marker_object` **40** in more detail. The `marker_object` **40** includes an `object_id` field **126** that holds a UUID, and a size field **128** specifies the size of the `marker_object` in bytes. A `marker_id` field **130** contains a UUID that identifies the marker data strategy, and a `num_entries` field **132** specifies the number of marker entries in the `marker_object` **40**. An `entry_alignment` field **134** identifies the byte alignment of the marker data, and a `name_len` field **136** specifies how many Unicode characters are held in the name field **138**, which holds the name of the `marker_object` **40**. Lastly, the `marker_data` field **140** holds the markers in a table. Each marker has an associated entry in the table.

FIG. 10B shows the format of a marker entry **141** such as found in the `marker_data` field **140**. An `offset` field **142** holds an offset in bytes from the start of packets in the `data_object` **47** indicating the position of the marker entry **141**. A time field **144** specifies a time stamp for the marker entry **141**. An `entry_len` field **146** specifies the size of an `entry_data` field **148**, which is an array holding the data for the marker entry.

The header section **28** may also include an `error_correction_object` **42** for an error correction method that is employed in the ASF stream. Up to four error correction methods may be defined for the ASF stream **16** and, thus, up to four `error_correction_objects` **42** may be stored within the header section **28** of the ASF stream **16**. FIG. 11 depicts the format of the `error_correction_object` **42**.

The `error_correction_object` **42** includes an `object_id` field **150** and a size field **152**, like those described above for the other objects in the header section **28**. The `error_correction_object` **42** also includes an `error_correction_id` **154** that holds UUID that identifies the error correcting methodology associated with the object **42**. The `error_correction_data_len` field **156** specifies the length in bytes of the `error_correction_data` entries **158** that hold octets for error correction. The `error_correction_object` **42** is used by the destination computer **12** (FIG. 1) in playing the ASF stream **16**.

FIG. 12 depicts a flowchart of how error correcting may be applied in the preferred embodiment of the present invention. In particular, an error correction methodology such as an N+1 parity scheme, is applied to one or more streams within the ASF stream 16 (step 160 in FIG. 12). Information regarding the error correcting methodology is then stored in the error\_correction\_object 42 within the header section 28 (step 162 in FIG. 12). The source computer then accesses the error correcting methodology information stored in the error\_correction\_object 42 in playing back the ASF stream 16 (step 164 in FIG. 12). Error correcting data is stored in the interleave\_packets 48.

The header section 28 of the ASF stream 16 may also hold a clock\_object 44 that defines properties for the timeline for which events are synchronized and against which multimedia objects are presented. FIG. 13 depicts the format of the clock\_object 44. An object\_ID field 166 holds a UUID to identify the object, and a size field 168 identifies the size of the clock\_object 44 in bytes. A packet\_clock\_type field 170 identifies the UUID of the clock\_type that is used by the object. A packet\_clock\_size field 172 identifies the clock size. A clock\_specific\_len field 174 identifies the size and bytes of the clock\_specific\_data field 176 which contains clock-specific data. The clock type alternatives include a clock that has a 32-bit source value and a 16-bit duration value, a clock type that has a 64-bit source value and a 32-bit duration value and a clock type that has a 64-bit source value and a 64-bit duration value.

The ASF stream 16 enables script commands to be embedded as a table in the script\_command\_object 45. This object 45 may be found in the header section 28 of the ASF stream 16. The script commands ride the ASF stream 16 to the client where they are grabbed by event handlers and executed. FIG. 14A illustrates the format of the script\_command\_object 45. Like many of the other objects in the header section 28, this object 45 may include an object\_ID field 178 for holding a UUID for the object and a size field 180 for holding the size in bytes of the object. A command\_ID field 182 identifies the structure of the command entry that is held within the object.

The num\_commands field 184 specifies the total number of script commands that are to be executed. The num\_types field 186 specifies the total number of different types of script\_command types that have been specified. The type\_names field 188 is an array of type\_names\_struct data structures. FIG. 14B depicts the format of this data structure 192. The type\_name\_len field 194 specifies the number of Unicode characters in the type\_names field 196, which is a Unicode string array holding names that specify script command types.

The command\_entry field 190 identifies what commands should be executed at which point in the timeline. The command\_entry field 190 is implemented as a table of script commands. Each command has an associated command\_entry element 198 as shown in FIG. 14C. Each such element 198 has a time field 200 that specifies when the script command is to be executed and a type field 202 that is an index into the type\_names array 196 that identifies the start of a Unicode string for the command type. A parameter field 204 holds a parameter value for the script command type.

The script commands may be of a URL type that causes a client browser to be executed to display an indicated URL. The script command may also be of a file name type that launches another ASF file to facilitate "continuous play" audio or video presentations. Those skilled in the art will appreciate that other types of script commands may also be used.

The header section 28 of the ASF stream 16 may also include a codec\_object 46. The codec\_object 46 provides a mechanism to embed information about a codec dependency that is needed to render the data stream by that codec. The codec object includes a list of codec types (e.g., ACM or ICM) and a descriptive name which enables the construction of a codec property page on the client. FIG. 15A depicts the format of a codec\_object 46. The object\_id field 206 holds a UUID for the codec\_object 46 and the size field 208 specifies the size of the object 46 in bytes. The codec\_ID field 210 holds a UUID that specifies the codec\_type used by the object. The codec\_entry\_len field 212 specifies the number of CodecEntry entries that are in the codec\_entry field 214. The codec\_entry field 214 contains codec-specific data and is an array of CodecEntry elements.

FIG. 15B depicts the format of a single CodecEntry element 216 as found in the codec\_entry field 214. A type field 218 specifies the type of codec. A name field 222 holds an array of Unicode characters that specifies the name of the codec and a name\_len field 220 specifies the number of Unicode characters in the name field. The description field 226 holds a description of the codec in Unicode characters and the description\_len field 224 specifies the number of Unicode characters held within the description field. The cbinfo field 230 holds an array of octets that identify the type of the codec and the cbinfo\_len field 228 holds the number of bytes in the cbinfo field 230.

As mentioned above, the data section 30 follows the header section 28 in the ASF stream 16. The data section includes a data\_object 47 and interleave\_packets 48. A data\_object 47 marks the beginning of the data section 30 and correlates the header section 28 with the data section 30. The packets 48 hold the data payloads for the media stream stored within the ASF stream 16.

FIG. 16 depicts the format of the data\_object 46. Like other objects in the ASF stream 16, data\_object 46 includes an object\_id field 232 and a size field 234. The data\_object 46 also includes a multimedia\_stream\_id field 236 that holds a UUID for the ASF stream 16. This value must match the value held in the multimedia\_stream\_id field 60 in the properties\_object 34 in the header section 28. The data\_object 46 also includes a num\_packets field 238 that specifies the number of interleave\_packets 48 in the data section 30. An alignment field 240 specifies the packing alignment within packets (e.g., byte alignment or word alignment), and the packet\_alignment field 242 specifies the packet packing alignment.

Each packet 48 has a format like that depicted in FIG. 17. Each packet 48 begins with an initial\_structure 244. The format of the initial\_structures 244 depends upon whether the first bit held within the structure is set or not. FIG. 18A depicts a first format of the initial\_structure 244 when the most significant bit is cleared (i.e., has a value of zero). The most significant bit is the error\_correction\_present flag 270 that specifies whether error correction information is present within the initial\_structure 244 or not. In this case, because the bit 270 is cleared, there is no error correction information contained within the initial\_structure 244. This bit indicates whether or not error correction is used within the packet. The two bits that constitute the packet\_len\_type field 272 specify the size of the packet\_len field 256, which will be described in more detail below. The next two bits constitute the padding\_len\_type field 274 and specify the length of the padding\_len field 260, which will also be discussed in more detail below. The next two bits constitute the sequence\_type field 276 and specify the size of the sequence field 258. The final bit is the multiple\_payloads\_

present flag **278** which specifies whether or not multiple payloads are present within the packet. A value of 1 indicates that multiple media stream samples (i.e., multiple payloads) are present within the packet.

FIG. 18B depicts the format of the initial\_structure **244** when the error\_correction\_present bit is set (i.e., has a value of 1). In this instance, the first byte of the initial\_structure **244** constitutes the ec\_flag field **280**. The first bit within the ec\_flag field is the error\_correction\_present bit **270**, which has been described above. The two bits that follow the error\_correction\_present bit **270** constitute the error\_correction\_len\_type field **284** and specify the size of the error\_correction\_data\_len field **290**. The next bit constitutes the opaque\_data flag **286** which specifies whether opaque data exists or not. The final four bits constitute the error\_correction\_data\_length field **288**. If the error\_correction\_len\_type field **284** has a value of "00" then the error\_correction\_data\_length field **288** holds the error\_correction\_data\_len value and the error\_correction\_data\_len field **290** does not exist. Otherwise this field **288** has a value of "0000." When the error\_correction\_data\_len field **290** is present, it specifies the number of bytes in the error\_correction\_data array **292**. The error\_correction\_data array **292** holds an array of bytes that contain the actual per-packet data required to implement the selected error correction method.

The initial\_structure **244** may also include opaque data **300** if the opaque\_data bit **286** is set. The initial structure includes a byte of flags **302**. The most significant bit is a reserved bit **304** that is set to a value of "0." The next two bits constitute the packet\_len\_type field **306** that indicate the size of the packet\_len field **256**. The next subsequent two bits constitute the padding\_len\_type field **272** that indicate the size of the padding\_len field **274**. These two bits are followed by another 2-bit field that constitutes the sequence\_type of field **276** that specifies the size of the sequence field **258**. The last bit is the multiple\_payloads\_present bit **278** that specifies whether are not multiple payloads are present.

The initial\_structure **244** is followed by a stream\_flag field **246** that holds a byte consisting of four 2-bit fields. The first two bits constitute a stream\_id\_type field **248** that specifies the size of the stream\_id field **314** within the payload\_struc **266**. The second most significant bits constitute the object\_id\_type field **250** and indicate the number of bits in the object\_id field **316** of the payload\_struc **266** as either 0-bits, 8-bits, 16-bits or 32-bits. The third most significant two bits constitute the offset\_type field **252**, which specifies the length of the offset field **318** within the payload\_struc **266** as either 0-bits, 8-bits, 16-bits or 32-bits. The least two significant bits constitute the replicated\_data\_type field **254** and these bits indicate the number of bits that are present for the replicated\_data\_len field **320** of the payload\_struc **266**.

The packet **48** also includes a packet\_len field **256** that specifies the packet length size. The sequence field **258** specifies the sequence number for the packet. The padding\_len field **260** contains a number that specifies the number of padding bytes that are present at the end of the packet to pad out the packet to a desirable size.

The packet **48** also contains a clock\_data field **262** that contains data representing time information. This data may include a clock license that contains a system clock reference that drives the progression of the time line under the timing model and a duration that specifies the effective duration of the clock license. The duration field limits the

validity of the license to a time specified in milliseconds. Under the model adopted by the preferred embodiment of the present invention, the source computer **10** issues a clock license to the destination computer **12** that allows the clock of the destination computer **12** to progress forward for a period of time. The progression of time is gated by the arrival of a new piece of data that contains a clock value with a valid clock license that is not expired.

The packet **48** also includes a payload\_flag field **264** that specifies a payload length type and a designation of the number of payloads present in the packet. The payload\_flag field **264** is followed by one or more payload\_strucs **266**. These structures contain payload information which will be described in more detail below. The final bits within the packet **48** may constitute padding **268**.

FIG. 19 depicts the payload\_struc **266** in more detail. The stream\_id field **314** is an optional field that identifies the stream type of the payload. The object\_id field **316** may be included to hold an object identifier. An offset field **318** may be included to specify an offset of the payload within the ASF stream. The offset represents the starting address within a zero-address-based media stream sample where the packet payload should be copied.

The payload\_struc **266** may also include a replicated\_data\_len field **320** that specifies the number of bytes of replicated data present in the replicated\_data field **322**. As was discussed above, for protection against possible errors, the packet **48** may include replicated data. This replicated data is stored within the replicated\_data field **322**.

The payload\_len field **323** specifies the number of payload bytes present in the payload held within the payload\_data field **325**. The payload\_data field **326** holds an array of payloads (i.e., the data).

The ASF stream may also include an index\_object **49** that holds index information regarding the ASF stream **16**. FIG. 20 depicts the format of the index\_object **49**. The index\_object includes a number of index entries. The index\_object **49** includes an object\_id field **324** and a size field **326**. In addition, the index\_object **49** includes an index\_id field **328** that holds a UUID for the index type. Multiple index\_name\_entries may be stored depending on the number of entries required to hold the characters of the name. For example, each entry may hold 16 characters in an illustrative embodiment.

The index\_object includes a time\_delta field **330** that specifies a time interval between index entries. The time represents a point on the timeline for the ASF stream **16**. A max\_packets field **332** specifies a maximum value for packet\_count fields, which will be described in more detail below. A num\_entries field **334** is a 32-bit unsigned integer that describes the maximum number of index entries that are defined within the index\_info array **336**. This array **336** is an array of index\_information structures. Each index\_info structure holds a packet field that holds a packet number associated with the index entry and a packet\_count field specifies the number of the packet to send with the index entry so as to associate the index entries with the packets. In FIG. 21, the index\_info array structure **336** holds N index\_information structures and each index\_information structure has a packet field **338A-338N** and a packet\_count field **340A-340N**.

While the present invention has been described with reference to a preferred embodiment thereof, those skilled in the art will appreciate that various changes in form and detail may be made without departing from the intended scope of the invention as defined in the appended claims. For

example, the present invention may be practiced with a stream format that differs from the format described above. The particulars described above are intended merely to be illustrative. The present invention may be practiced with stream formats that include only a subset of the above-described fields or include additional fields that differ from those described above. Moreover, the length of the values held within the fields and the organization of the structures described above are not intended to limit the scope of the present invention.

We claim:

1. In a computer system operable for connecting to a transport medium, a method of encapsulating multiple streams of data into an aggregated data stream to be transmitted on the transport medium comprising:

determining a packet size for the transport medium;

storing, on a storage device, at least one packet containing information about the aggregated data stream to form a header section in a logical structure that defines the aggregated data stream;

storing, on the storage device, packets containing samples of data from the multiple data streams to form a data section in the logical structure;

designating a portion of at least one packet in the data section for holding error correcting data; and

storing, on the storage device, error correcting data in the designated portion of the at least one packet in the data section, wherein the error correcting data identifies a corresponding error correcting method for each of the encapsulated data streams, and further wherein the aggregated data stream is stored on the storage device prior to receiving a request for transmission of the aggregated data stream on the transport medium from a destination computer.

2. The method of claim 1, further comprising using the error correcting data stored in the designated portion of at the least one packet to correct an error when the aggregated data stream is received by the destination computer on the transport medium.

3. The method of claim 1 wherein the error correcting data is stored in multiple packets in the data section.

4. The method of claim 1 wherein the error correcting data holds parity bits.

5. The method of claim 1 wherein packets containing samples of data from a first of the multiple streams hold a different type of error correcting data than packets containing samples of data from a second of the multiple streams.

6. The method of claim 1 wherein the information in the header section of the logical structure indicates what error correcting methodology is used with the at least one packet in the data section that holds error correcting data.

7. The method of claim 1, further comprising:

transferring the packets from the storage medium as arranged in the logical structure across the transport medium to the destination computer.

8. The method of claim 1 wherein at least two of the multiple streams of data are of different media.

9. In a computer system operable for connecting to a transport medium, a method of building an aggregated data stream in a logical structure on a storage medium in the computer system comprising:

determining a packet size for the transport medium;

storing data from multiple streams of data into packets;

storing error correcting data in at least some of the packets;

encapsulating the packets into the logical structure for the aggregated data stream; and

storing information regarding a plurality of error correcting methods that are employed for the packets in the aggregated data stream into the logical structure,

wherein the aggregated data stream is built prior to receiving a request for transmission of the aggregated data stream on the transport medium from a destination computer.

10. The method of claim 9 wherein the logical structure for the aggregated stream includes a header section and the information regarding the error correcting methods that are employed is stored in the header section.

11. The method of claim 10 wherein multiple error correcting methods are employed.

12. The method of claim 11 wherein a separate object is stored in the header section for each error correcting method to encapsulate the information regarding the error correcting method.

13. In a computer operable for connecting to a transport medium, a computer-readable storage medium holding instructions for encapsulating multiple streams of data into an aggregated data stream to be transmitted on the transport medium:

determining a packet size for the transport medium;

storing, on a storage device, at least one packet containing information about the aggregated data stream to form a header section in a logical structure that defines the aggregated data stream;

storing, on the storage device, packets containing samples of data from the multiple data streams to form a data section in the logical structure;

designating a portion of at least one packet in the data section for holding error correcting data; and

storing, on the storage device, error correcting data in the designated portion of the at least one packet in the data section, wherein the error correcting data identifies a corresponding error correcting method for each of the encapsulated data streams, and further wherein the instructions are executed prior to receiving a request for transmission of the aggregated data stream on the transport medium from a destination computer.

14. The computer-readable storage medium of claim 13 further holding instructions that store the error correcting data in multiple packets in the data section of the logical structure.

15. The computer-readable storage medium of claim 13 wherein the medium holds instructions for encapsulating a first type of error correcting data for packets that contain samples from a first of the streams of data and encapsulating a second type of error correcting data for packets that contain samples from a second of the streams of data.

16. In a computer system, a computer-readable storage medium holding a logical structure that encapsulates multiple streams of data, the logical structure comprising:

packets of data from the multiple streams of data forming a data section of the logical structure and stored on mass storage for subsequent transmission over a transport medium when requested by a destination computer; and

error correcting data within at least some of the packets at designated locations, wherein the error correcting data identifies a plurality of error correcting methods.

13

17. The computer-readable storage medium of claim 16 wherein the logical structure further comprises a header section in which information regarding an error correcting method that uses the error correcting data is stored.

18. The computer-readable storage medium of claim 17 5 wherein the header section holds information regarding multiple error correcting methods.

19. The computer-readable storage medium of claim 16 wherein the multiple streams of data include at least two streams of different media. 10

20. In a computer system, a computer-readable storage medium holding instructions for:

receiving a logical structure that holds multiple streams of data, wherein the streams of data include samples that are stored in packets in the logical structure and wherein at least some of the packets include error correcting data that was stored in the packets prior to a request being transmitted by the computer system that caused the logical structure to be received, wherein the 15

14

error correcting data identifies a plurality of error correcting methods; and

extracting the error correcting data from at least some of the packets as needed to correct errors according to the identified error correcting methods.

21. A data processing system having:

a source computer with a storage;

a logical structure stored in the storage for encapsulating multiple data streams into an aggregated data stream defined by the logical structure, data from the data streams being incorporated in packets prior to a request being received by the source computer to transmit the aggregated data stream to a destination computer; and error correcting data encapsulated in the packets, wherein the error correcting data identifies a corresponding error correcting method for each of the encapsulated data streams.

\* \* \* \* \*